

Detecting Lateral Spear Phishing Attacks in Organizations

A Dissertation

Submitted by

Aniket Bhadane

MIS: 121622001

in partial fulfilment for the award of the degree


of

M.Tech. Computer Engineering

Under the guidance of

Dr. Sunil B. Mane

College of Engineering, Pune



**DEPARTMENT OF COMPUTER ENGINEERING AND
INFORMATION TECHNOLOGY,
COLLEGE OF ENGINEERING, PUNE - 5**

June, 2018

**DEPARTMENT OF COMPUTER ENGINEERING AND
INFORMATION TECHNOLOGY,
COLLEGE OF ENGINEERING, PUNE**

CERTIFICATE

Certified that this dissertation, titled
“Detecting Lateral Spear Phishing Attacks in Organizations”
has been successfully completed by

Aniket Bhadane

MIS: 121622001

and is approved for the partial fulfillment of the requirements for the degree of
“M.Tech. Computer Engineering”.

SIGNATURE

Dr. Sunil B. Mane

Dissertation Guide

**Department of Computer Engineering
and Information Technology,
College of Engineering Pune,
Shivajinagar, Pune - 5.**

SIGNATURE

Head

**Department of Computer Engineering
and Information Technology,
College of Engineering Pune,
Shivajinagar, Pune - 5.**

Acknowledgments

I would like to express my sincere and deepest gratitude to my dissertation guide Dr. Sunil B. Mane for his continuous support and understanding throughout this journey. I am grateful to all the volunteers who shared their email history for preparation of the dataset without which the research would have seriously hampered. I would like to thank all the faculty members of Department of Computer Engineering & IT and the Helpdesk Team for providing me with the necessary help whenever required. I express my gratitude to my colleagues with whom I had valuable discussions on the project. Lastly, but importantly, to my parents who have loved me unconditionally all the time.

Abstract

Social engineering used by attackers is psychological manipulation of people to make them perform certain actions of attackers' benefit. Phishing attacks masquerade as a trustworthy source, often spread using socially engineered messages using media such as emails, social media, SMS, online multiplayer games, VoIP, etc. The socially engineered messages persuade users to click on a malicious/phishing link in an email; enter sensitive information on a phishing website; perform actions such as money transfer, installing malicious software etc. Phishing continues to be a problem for organizations and individuals by causing huge tangible and intangible losses.

Lateral Spear Phishing attack is a powerful class of social engineering attack carried out using compromised email account(s) within the target organization. Spear phishing attacks are difficult to detect due to the nature of these attacks. The inclusion of lateral attack vector makes detection more challenging. We present an approach to detect Lateral Spear Phishing attacks in organizations in real-time. Our approach uses features derived from domain knowledge and analysis of characteristics pertaining to such attacks, combined with our scoring technique which works on non-labelled dataset. We evaluate our approach on several years' worth of real-world email dataset collected from volunteers in our institute. We were able to achieve false positive rate of below 1%, and also detected two instances of compromised accounts which were not known earlier. Comparison of our scoring technique with machine learning based anomaly detection techniques shows our technique to be more suited for practical use. The proposed approach is primarily aimed to complement existing detection techniques on email servers. However, we also developed a Chrome browser extension to demonstrate that such system can also be used independently by organizations within their network.

Contents

List of Tables	ii
List of Figures	iii
1 Introduction	1
2 Existing State of Research	4
2.1 Overview of Literature on Phishing	4
2.2 Related Work	9
3 Concepts and Challenges	11
4 Data Characterization	13
5 Design of Scoring System	16
5.1 Feature Characterization	16
5.1.1 Behavioral Context Feature Category	16
5.1.2 IP Feature Category	18
5.1.3 Display Name Feature Category	20
5.1.4 FQDN Feature Category	20
5.2 LSP Scoring	22
6 Evaluation	25
6.1 Comparison with Anomaly Detection Techniques	27
7 Discussion	29
7.1 Limitations	29
7.2 Chrome Extension - COEP Kumpan	30
8 Conclusion	32
Bibliography	40
List of Publications	41

List of Tables

4.1	Number of Emails in our Dataset	13
4.2	Attack instances with Phishing URLs and Display name spoof	13
5.1	Summary of features fed into our Scoring system	22
6.1	Comparison of performance of our Scoring Technique with that of Anomaly Detection Techniques	26

List of Figures

2.1	Classification/Taxonomy of Anti-Phishing Techniques.	6
4.1	Number of unique emails in our dataset with intervals of 4 months. . . .	14
4.2	Number of attack emails in our dataset with intervals of 4 months.	14
5.1	Example of a hierarchy related to computer department in an institute. Note: The figure is not representative of any hierarchy and is meant only for illustration purpose to explain the concept of hierarchy. Organizations can decide the format in which they want to represent their hierarchy. . .	17
5.2	Complementary cumulative distribution function (CCDF) of number of times a FQDN was seen in our dataset.	21
6.1	CCDF of LSP scores for benign and lateral spear phishing emails given by the LSP Scoring system.	27
7.1	(a) shows an example of the UI element in an email in Outlook Web. (b)- (d) show examples of UI element with messages, either nudged or on mouse hover based on user preferences.	31

Chapter 1

Introduction

Phishing is a semantic attack. It exploits human vulnerabilities by targeting the way humans interpret content displayed on the system. This is commonly done using social engineering. A phishing attack or a phishing taxonomy using phishing messages has three major stages: First stage corresponds to Attack Lure where the victim(s) receive a socially engineered phishing message with the sender masquerading/impersonating as a trusted identity. Second and Third stages correspond to Attack Exploit, where victim(s) perform action as suggested in the message, and attacker exploits the action performed by the victim(s) for his benefit.

Spear phishing attacks have been in limelight of cyber-attacks, costing organizations and individuals' huge tangible and intangible losses [27]. Lazarus group known for its involvement in WannaCry ransomware attack, recently launched a malicious spear phishing campaign with lure as a job opening for CFO role at a cryptocurrency company [51]. With sharp increase in prices of cryptocurrencies, crypto-phishing has seen steep increase since 2017 [47]. Also recently, sensitive data of upto 30,000 Medicaid recipients was compromised when an employee of Florida's Agency for Health Care Administration fell for a phishing attack. Organizations, be it government or corporate or non-profit, have been compromised by these types of attacks [26]. Such emails also take advantage of current events such as Equifax hack [42].

Spear phishing emails are targeted towards specific individuals or organizations and trick target(s) using social engineering techniques into performing actions that are of benefit to the attacker. The actions could be clicking on a URL in the email and entering credentials on the resulting phishing page; revealing sensitive information about self or organization; doing wire transfer; etc. Adversaries can easily launch spear phishing campaigns with availability of many toolkits to aid the tasks. Phishing-as-a-service [37] has made con-

ducting attacks even more convenient for them. Adversaries also use various mechanisms for extending the lifespan of phishing sites by limiting the visibility of the attack [52]. Defending from such attacks is difficult as they are crafted to appear legitimate. 70% of the spear phishing emails are opened by their targets and 50% of these users are found to click on URLs in the email [48, 27].

A lateral attack involves compromising a node in the target organization, which gives attackers a foothold in the network. Attackers can then move laterally within the organization targeting sensitive data such as high level personnel credentials, financial data, intellectual properties, etc. Lateral Spear Phishing is a powerful class of attack in which spear phishing emails are sent to people in the organization from a compromised email account within the organization itself. Such types of attack are particularly rare than other types of spear phishing attacks and can prove to be highly expensive for organizations. There is no generic solution for detecting these types of attacks and solutions need to incorporate the domain knowledge of their problem space to get best results [35].

We collected several years of emails from volunteers in College of Engineering Pune (COEP), an autonomous engineering institute in the state of Maharashtra in India. The institute has faced various lateral spear phishing attacks within last few years. COEP uses corporate Outlook as its email service provider. The lateral spear phishing attacks faced by COEP are primarily credential seeking or trying to establish a conversation to extract sensitive information. We made effort to collect emails such that the dataset contains all known attack instances.

Our goal is to create a practical and deployable technique which can detect fair number of lateral spear phishing attacks in real-time, keeping low false positive rate. Standard machine learning techniques suffer due to the high imbalance of dataset and have their false positive rate impractical for use by organizations. Importantly, intra-domain emails (emails to and from the same domain) do not tend to get blocked as they have a much higher ‘*trust*’ score [68]. We derive a compact set of features from the analysis of lateral spear phishing attacks. Further, our scoring technique leverages domain knowledge of the attacks on the organization. This also implies that administrators can customize the technique to integrate their domain knowledge pertaining to their organization. We compare our technique with machine learning based anomaly detection techniques and show that our technique gives better accuracy and false positive rate than the compared machine learning techniques.

Our technique is mainly aimed to be implemented on email servers providing service to organizations, to complement existing detection techniques. If that is not possible due to plausible reasons, an organization can implement such technique within their network,

with the help of plugins, to aid users in decision-making. We also developed a Chrome Extension which uses the API provided by our scoring system, to aid COEP's users in taking better informed decisions when dealing with emails on COEP's Outlook email domain. The major contributions of this work are:

- Usable real-time approach for detection of lateral spear phishing attacks.
- Use of domain knowledge and characteristic features of these attacks, alongwith our scoring technique which requires non-labelled dataset.
- Suitable for use by organizations where intra-domain emails have much higher trust score, and aggressive heuristics are not used to avoid false positives and false negatives.
- System which is primarily aimed to be implemented on email servers complementing existing detection mechanisms, but can also be implemented independently by organizations if they don't manage their email server.

In Chapter 2, we present an overview on the existing literature on phishing and the literature related to our work. We discuss about some concepts surrounding these attacks and challenges in detection of such attacks in Chapter 3. In Chapter 4, we provide the characterization of the dataset that we collected. We describe the features we use and our scoring technique in Chapter 5. In Chapter 6, we present the evaluation of our scoring technique and compare its performance with anomaly detection techniques. In Chapter 7, we provide discussion, limitations of the system, and some details on the Chrome Extension we developed for use by COEP. Chapter 8 concludes the dissertation.

Chapter 2

Existing State of Research

2.1 Overview of Literature on Phishing

Zhang et al. [81] had categorized the past work on phishing into four categories: understanding why people fall for phishing, automated software techniques to detect phishing, training people to not fall for phishing attacks, and better user interfaces to help people make better decisions when confronted with an attack. However, we consider that these can be generalized into two main categories: one which try to understand human psychology involved in phishing, and other which develop techniques to mitigate phishing.

Phishing Psychology

Phishing messages generally use sentiments and psychological principles of influence - authority, scarcity, curiosity, fear, urgency, social proof etc [56]. [18, 79] found that even experienced users can fall prey to visual deception attacks. Also, [21, 20] showed that traditional ways of user awareness did not contribute much in reducing their susceptibility to attacks. Canfield et al. [9] conducted experiments to find performance of users in Detection - deciding whether an email is phishing or not, and Behavior - deciding what action to take on a mail. With the help of signal detection theory, they show that phishing-related decisions are sensitive to users' confidence and view on consequences. Also, [73] showed that judgmental confidence has more effect on users' behaviors than the confidence due to people's beliefs about their capabilities. The research findings of [57] show that younger users were found to be most susceptible to scarcity (e.g.: "once in a lifetime opportunity" type messages), while older users were most susceptible to reciprocation (e.g.: luring user to install malware by offering a free gift). [14] found that the variation

in workload and the number of unimportant mails that an employee interacts with on a daily basis, are related with users' phishing susceptibilities. Vishwanath [72] studied users' Facebook usage habits and their relation with user susceptibility to social media phishing attacks. His results show that habitual facebook use, resulting in automaticity of response is a leading factor involved in an individual's victimization in social media attacks.

Summarizing, people of different demographics have different psychological vulnerabilities. People have different habits and perceptions in different environments, with habituation playing a major role in user susceptibility to social media phishing attacks. Users' confidence and their view on consequences are also related to users' susceptibility to phishing attacks.

Anti-Phishing Techniques

From our understanding of literature, the tree classification/taxonomy of anti-phishing techniques can be visualized as shown in Figure 2.1. Solutions to mitigate phishing can be briefly classified into: Make it Invisible; Train Users; Better Interfaces.

Heuristics to detect phishing emails at server side mainly use machine learning techniques. Researchers have investigated large number of features which can be used to identify phishing emails [3, 32]. These features can be categorized as: Structural, Link, Element, Spam filter, Semantic Keywords, and Dynamic Markov Chain Features. Coupled with these categories of features, researchers have explored various classes of machine learning algorithms, viz. Bag-of-Words models, Multi-Classifiers, Clustering, Classifiers model based, Multi-layered, and Evolving Connectionist System [3, 34, 32]. Detecting phishing attacks at network level is based on domain and IP address blacklisting, such as DNSBL and Snort. Some researchers have also proposed systems which try to identify whether an email is sent by the intended sender, by constructing profiles of the senders [41, 70, 22]. It can be observed from the existing literature, that there does not exist a standard classifier for detection of phishing messages and no heuristic technique proposed till now is foolproof. Most classifiers are based on supervised learning and most of the work is done offline. Although heuristics are able to detect some of the attacks at zero-hour, they also have side-effects of false positives. False positives are a major liability when deploying aggressive technical solutions in corporate settings. A business related legitimate message being flagged as phishing can have serious implications on the business.

Email Authentication and Verification techniques (SPF, DKIM, DMARC) provide the framework for businesses to combat spoofed emails. They are used to combat phishing

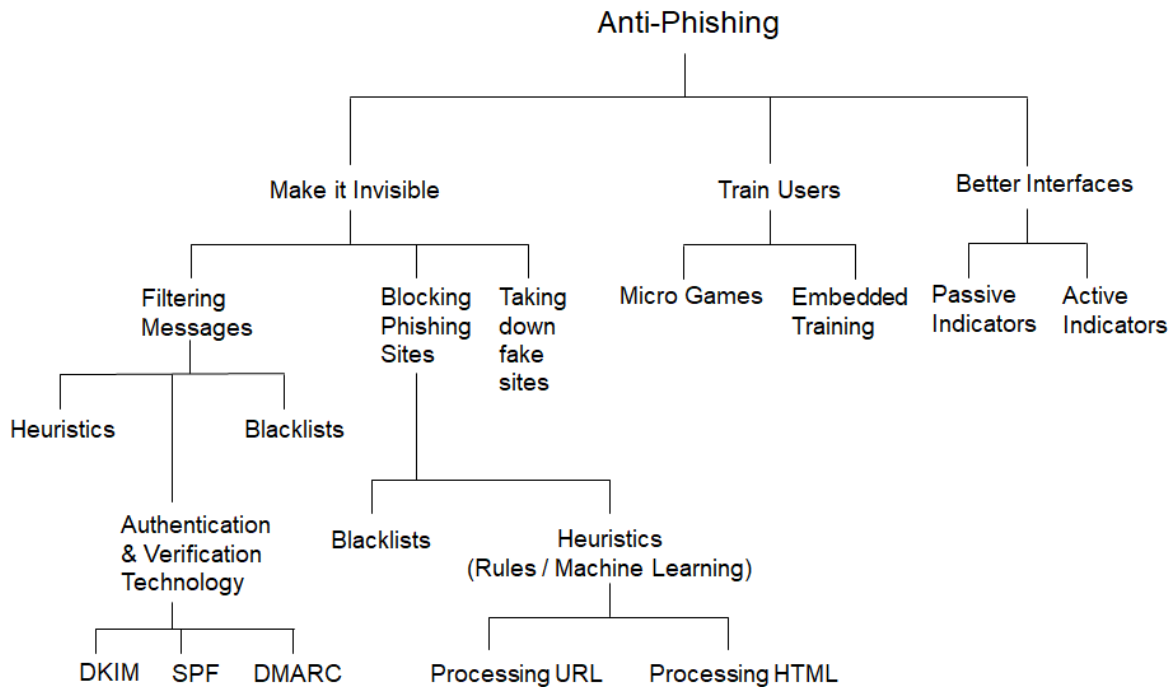


Figure 2.1: Classification/Taxonomy of Anti-Phishing Techniques.

techniques using forged sender addresses that appear to originate from legitimate organizations. SPF informs the receivers which mail servers are used to send mail from sender's domain. DKIM authenticates the source and its contents using sender's signature. DMARC builds upon SPF and DKIM, and defines a scalable mechanism by which a mail sender can express domain level policies governing how messages claiming to come from his or her domain should be handled when they do not fully align with DKIM and SPF lookup results. However, these suffer many shortcomings and 2017's DMARC adoption report by Agari mentions that two-thirds of the Fortune 500 are yet to implement DMARC [1]. Although, these techniques can prevent against phishing attacks using email address spoofing, but do not provide protection against phishing attacks which do not use email address spoofing.

DNS-based Blackhole Lists (DNSBLs) can be used by mail servers to get suggestion on an email's origin. Some DNSBLs, such as Spamhaus DBL, maintain list of domains used for sending phishing emails and claim almost zero false positives, so that it is low-risk to be used by production mail systems to reject emails that are flagged by it. There have been criticisms against DNSBLs such as blocking legitimate emails from shared mail servers, treating IP address from ISPs as dynamic and hence not suitable for sending mails directly, etc. [19]

A new website can be classified as phishing or not by analyzing its URL and/or HTML

content. URL features can be categorized as Lexical-based and Host-based [50]. Lexical features are the textual properties of the URL. Host-based features include where the site is hosted, who owns the site, and how is it managed. These alongwith HTML content based features are used to define rules or train machine learning classifiers to detect phishing. Numerous studies [62, 13, 33, 76] use various features extracted from URL and HTML of websites to determine whether they are phishing or not. [39] provided analysis of Visual Similarity based approaches for phishing websites' detection. [75, 63] look at the hyperlink relationships between phishing sites and their targets. In 2017, Cui et al. [15] monitored a total of 19,066 phishing attacks over a period of ten months and found that over 90% of these attacks were actually replicas or variations of other attacks in the database. Their clustering results show that a small group of attackers could be behind a large part of the current attacks. Studies have also used Whitelists to maintain characteristics of trusted websites and then compare them with a new website to see if its phishing [10, 5].

A phishing websites' blacklist is a list of websites which have been marked as phishing by the particular blacklist service. The most popular of these are Google Safe Browsing, PhishTank, and APWG. Browsers such as Google Chrome, Safari, Firefox, Opera, and Vivaldi are known to use Google Safe Browsing blacklist [77]. PhishTank is a community where users can submit phishing websites, which are then verified using community voting, and added to their blacklist of phishing websites. The blacklist is used by organizations and services like Yahoo Mail, McAfee, Kaspersky, APWG, Opera and many more [61]. The Anti-Phishing Working Group (APWG) is an international group which prepares and distributes phishing reports to its paying members [60]. [67] found that blacklists were ineffective when protecting users against attacks at hour zero, were updated at different speeds and extent.

Once a phishing page is detected, it must be reported to the respective ISP/hosting company in order to take down such content. Several organizations try to report such sites to ISPs. The time required to take down a site depends on the service provider [64]. In an approach called Fast Flux, attackers try to hide a phishing website behind a large number of compromised proxies. [36, 15] observed that even though a phishing website is taken down, the attackers launch the website from different domains, IPs and URLs, without much modifying the website. They also mention that certain attacker groups are responsible for large number of attacks, and use same resources for carrying out different attacks. Identifying these groups can stop a large number of attacks.

User Education and Training is an important part of online security, but it has not received enough attention from the phishing community. Some studies have shown training

programs to be helpful [45, 66], while some disagree or say that training has mixed benefits [11, 30]. Our learning from the literature was that the best way of training against phishing attacks is “at the moment of attack - periodic - demographic tailored” training [46, 45, 57]. Research on User Training against phishing can be categorized into developing Micro Games or Embedded Training.

Anti-Phishing Phil [69] is a micro game which educates users about various parts and cues to identify phishing and legitimate URLs. They used leaning sciences to build intervention based designs, and found it to be more effective than security notices emailed by companies to users. Control-Alt-Hack [16] lets users interact with cards in a board game, through which it teaches users about different social engineering techniques attackers use. What.Hack [74] presents a sequence of puzzles in a story-based game context, to teach users about social engineering threats. Some of the commercial offerings available are Email Security or Anti-Phishing Phyllis by Wombat Security [78]. In Embedded Training, the educational material is embedded, i.e. integrated into the daily primary tasks of users. Few organizations, for example Indiana University [38], have used an approach where simulated phishing emails are sent to users to train them and the results showed an improvement in users’ phishing detection ability. In PhishGuru [44], simulated phishing emails are sent to users in their normal use of emails, and uses interventional educational messages to educate them. When the users fall victim to the attack, such as clicking on a link in the phishing email, the system teaches users about the attack.

Well maintained organizational policies help in dealing with potential exploits occurring, without relying on the technical capabilities of users [34]. Researchers have proposed anti-phishing frameworks with such policies [28]. Also some internationally standardized frameworks have integrated such policies [2] and even few Governments have proposed guidelines [55].

One line of research aims to have better design of phishing warnings in browsers or tools or services so that users can take better informed decisions against such attacks. Our learning from the literature is that warnings should interrupt users at the right time, with precise message, and avoid habituation [36, 79, 65, 24]. A passive indicator does not interrupt the user. These indicators contain text and/or icons of some specific colors. For example, browser security indicators. Research has shown that many users don’t look at or ignore passive indicators [79]. [25] proposed three security indicators considering design constraints such as scale, shape, interpretation. These are now used with Google Chrome 53. [8] designed “Inhibitive Attractors”, which are UI elements that draw user’s attention to important part of the message in a security dialog to help user make better informed decision. Active indicators interrupt user’s activities. [54] proposed use of social saliency

nudges in emails, for users to take better informed decisions in case of phishing emails. [24, 79] showed active warnings to be more effective than passive warnings. Users pay attention to warnings only when forced to do so and using a user-understandable warning message. Also, results from the active warning indicators depends on their design and implementation, as noted that warnings in one browser gave better results than other browser. Interruptions have shown to increase effectiveness of warning messages, but they should not be too frequent such that the user gets habituated to them and starts ignoring them.

2.2 Related Work

The closest to our work is [35], where the authors proposed an approach for detecting credential spear phishing attacks in settings where attack instances are as few as one to millions of benign emails. Their work draws on the SMTP, NIDS and LDAP logs of 370 million anonymized emails from their organization. They rank top n suspicious events every nightly for an analyst to review. The n is their daily budget of alerts that the analyst has to review (budget of 10 alerts mentioned in the paper). They detect attacks of three impersonation types with three different sub-detectors. And each sub-detector detecting attacks at two different stages of attack (lure and exploit). At their daily budget, they achieved very low false positive rate of 0.005%. However, since they do not mean to alert/detect all the attacks in the given day, the actual false positive rate in this case is not known. Our system would enhance capabilities of such detector, for detecting lateral spear phishing attacks in particular.

IdentityMailer [70] builds historical profiles for users based on their habits of email writing, composition and recipient interaction. Working at sender's side, it compares the sending user's historical profile with the current email being sent. Similar to this, Email-Profiler [22] works on the recipient's side and builds historical profiles from incoming emails. They use email's metadata and stylometric features with a SVM classifier. Given the organization's conditions on privacy and confidentiality, such techniques could be of aid to our system by considering textual content in the email bodies.

Pecchia et al. [59] worked on detecting compromise of account based on the activities from the account post-login. They used various security indicators such as suspicious command-line activity, file downloads, etc. to train a naive bayesian network. Although they detected all attacks in their dataset, it also cost them huge number of false positives in doing so. The authors suggest that the approach is more suited to guide the administrators in dealing with potential compromises.

Freeman et al. [29] developed a statistical Bayesian framework to detect suspicious login attempts. They propose use of what they call reinforced authentication, which checks for additional information beyond credentials during login attempts. They extract this information from the HTTP session logs of authentication sessions. Their results show false positive rate of 10% which they reason as a business decision. This false positive rate is quite high for other organizational settings. Also working on authentication logs, Zhang et al. [80] employed logistic regression classifier on features extracted from the authentication logs in their academic institution.

Laszka et al. [49] worked on the problem of strategic email threshold selection for filtering spear phishing emails using concepts of game theory. The authors consider the case where the attacker selectively targets users considering various factors which maximize the value gained. They theoretically prove that defenders may reach Nash equilibrium on the long term. However, they do not prove the performance of such system on any dataset.

On the grounds of social networks, [23, 71] conducted large scale experiments in which they construct clusters based on content and URL similarity of Facebook posts / Twitter tweets. [23] looks for deviation in the users' behavior as opposed to their historical profiles. Whereas [71] use a logistic regression classifier to categorize the clusters for various types of accounts. These studies show that account compromise on social networks is largely due to phishing and the compromised accounts are in turn used to conduct large scale phishing campaigns. Dewan et al. [17] trained four machine learning classifiers on the stylometric features of emails and social features from LinkedIn profiles. However, their results show that their chosen features from LinkedIn did not help in identifying spear phishing emails.

Chapter 3

Concepts and Challenges

Lateral spear phishing attack first involves compromise of an email account within the target organization. This serves as an entry point and foothold in the organization's network for an attacker. The compromise can be due to several reasons such as brute force, theft, etc. On setting the foothold, the attacker can move laterally within the organization seeking sensitive information. The attacker can also choose to stay passive, with the objective of just scanning the email conversations of the victim to get further sensitive information. The presence of the attacker within the organization will remain till any suspicious activity is detected. Although remote access is essential for workflow of organizations, it also provides chances for the attackers to get into the organization's network. Stolen credentials are also sold on dark web from an average of \$8 to maximum of \$190 each [43]. Zhang et al. [80] also found that compromised credentials of university accounts were used to access scholarly articles and bypass censorship.

Email account compromise can be detected at two stages:

1. during login; by analyzing properties of login activity.
2. post-login; by analyzing email sending activity.

Detection approaches at both stages have their pros and cons, and need to complement each other. Our detection approach detects attacks at the second stage. We look for suspicious email sending activity to detect account compromise regardless of the way of compromise. Note that we refer to lateral spear phishing attacks as those which are conducted from compromised account within the organization, not from user's personal email account.

One challenge in dealing with these attacks is that such incidences occur infrequently, which makes this problem similar to finding needle in haystack. But a successful com-

promise from these attacks could be devastating for the organization. Such infrequent occurrence of attacks leads to base-rate issues with imbalanced dataset having small training set. Due to these issues, success of standard machine learning techniques is improbable. Other issues include limited history of users and benign churn in header values, which further degrades the results of standard machine learning techniques [35]. We provide comparison of our technique with machine learning algorithms and discuss more about their results in 6.1.

Another challenge is that intra-domain emails have higher trust score and hence an anomalous intra-domain email is not filtered [68]. Phishing emails which should normally get detected when sent from a personal email account, do not get detected when sent from an email account within the domain, mainly due to the higher trust score. Aggressive heuristics are not used largely as they tend to have higher false positives and false negatives for organizational settings [58] which hamper productivity of the organization. Legitimate employee-to-employee emails can look a lot like spam (e.g.: got my discount, follow this referral link [URL]). This is issue of ease of use. The sensitivity level of intra-domain filters can be modified by the administrators in the organization. But it involves risk of false positives and needs sufficient testing and to actually give better results without hampering normal flow of business.

Multi-factor authentication can increase cost to attackers for infiltrating the network. However, these authentication techniques suffer from deployability and usability issues. The success of these techniques depends largely on its implementation and is not difficult for attackers to get past weaker implementations [7]. Organizations including universities and national laboratories find it difficult to adopt these authentication techniques in practice [40].

Lateral spear phishing attack can also have avalanche effect within the organization which further amplifies the damage caused by the attack. Taking example of one of our case, a staff's compromised account led to compromise of multiple students' accounts using spear phishing emails. These accounts in turn were used to conduct further attacks and compromise more number of accounts. Further, in the current scenario, organizations largely depend on users to report spear phishing attacks [35]. In the next chapter, we discuss about our dataset and provide its characterizing details.

Chapter 4

Data Characterization

Table 4.1: Number of Emails in our Dataset

Total number of emails	113,726
Unique emails	19,087
Unique intra-domain emails	12,064

Table 4.2: Attack instances with Phishing URLs and Display name spoof

Attack instances containing phishing URLs	22/27
Attack instances involving display name spoof	04/27

Our work draws upon on the emails collected from volunteers in COEP. COEP uses Office365 Outlook for its emailing needs. Due to privacy and confidentiality reasons, we do not have access to the entire emailing history of our organization. However, we collected emails from 40 volunteers in the organization. The volunteers' emails were downloaded and converted to a processable format. The emails contain both their respective header and body. The dataset contains 3.5 years' worth of emails, ranging from July 2014 to January 2018. The total number of emails we collected from 40 volunteers is 113,726. On merging the emails from individual volunteers, our dataset had 19,087 unique emails. Out of these, 12064 are intra-domain emails. This is recapped in Table 4.1.

We made effort to have all known attack instances in the dataset even though we do not have all the benign emails of the organization. The dataset contains entire history of emails of student volunteers from various academic years (First Year BTech, Second Year BTech, etc.) and also from some of the staff members. The types of spear phishing attacks that the attacker(s) carried out after compromising an account in the organization were either to steal credentials using a phishing site impersonating COEP's site, or to establish

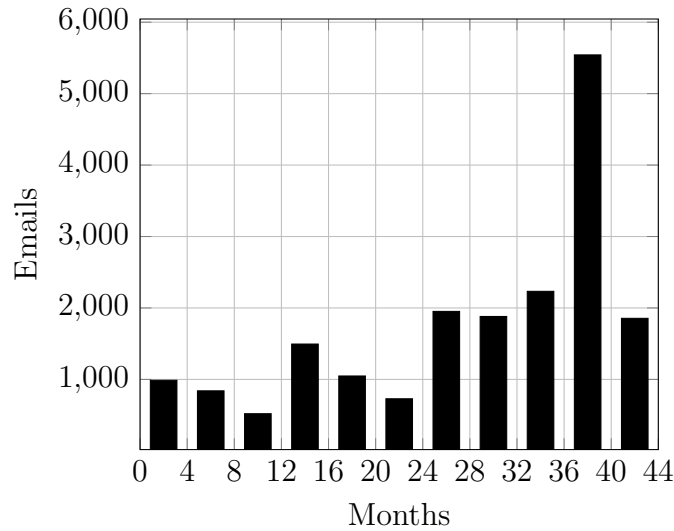


Figure 4.1: Number of unique emails in our dataset with intervals of 4 months.




Figure 4.2: Number of attack emails in our dataset with intervals of 4 months.

communication with another target within the organization to get sensitive information.

Our dataset contains total of 27 account compromise instances, 25 of which were known earlier and 2 were detected after we test our system with the dataset. The 27 compromise instances correspond to 323 spear phishing emails. Once an account was compromised, it was used to send spear phishing emails to other targets in the organization. The compromised account access was revoked only after users reported the attack to the administration team. Since the administration team is not aware of any ongoing attack, it requires reporting by the users. Some of these spear phishing emails were even sent in bulk to others in the organization impersonating as helpdesk of the organization. It is not known how many users have responded to the spear phishing emails given the large

number of users that the spear phishing emails were sent to.

Table 4.2 shows the number of attack instances that contained phishing URLs and the number of attack instances which involved display name spoof. All the phishing websites were impersonating as COEP's site to steal target's credentials and also all display name spoof instances involved impersonation of someone within the organization.

Our dataset contains more number of emails towards the end of the observation period as the dataset contains emails from all volunteers in that period; which is not necessarily true at the start of the observation period since we have volunteers from different academic years. Figure 4.1 shows plot of number of emails in our dataset with intervals of 4 months. Also, Figure 4.2 shows the plot of number of attack emails with intervals of 4 months. The number of such attacks increased tremendously from the latter half of 2016.

Chapter 5

Design of Scoring System

The main goal of the scoring technique is to have low false positive rate and detecting fair number of true positives. Another goal is to be real-time in order to process emails immediately as they are received. We leverage the structured information in email headers and one element in email body. The scoring technique determines a score for every received email and decides actions to take correspondingly. The following two sections describe about the features that are feed into our scoring technique, and then the scoring technique itself.

5.1 Feature Characterization

We derive Four Feature Categories from the characteristics of the lateral spear phishing emails we observed in our organization. Each feature category contains one or more features. Based on its nature, a feature could be either context-based or history-based. A context-based feature analyzes a specific property/context of the email, whereas history-based feature compares itself with a historically learned profile.

5.1.1 Behavioral Context Feature Category

Once an account within the organization is compromised, the attacker can send spear phishing emails to other users in the organization who are not related to the original user of compromised account. We try to identify this behavior. A natural solution is to find frequent email sending patterns from the email history. However, this does not suit to our case as we explain later in this section. The idea is to check whether there is a ‘*meaningful*’

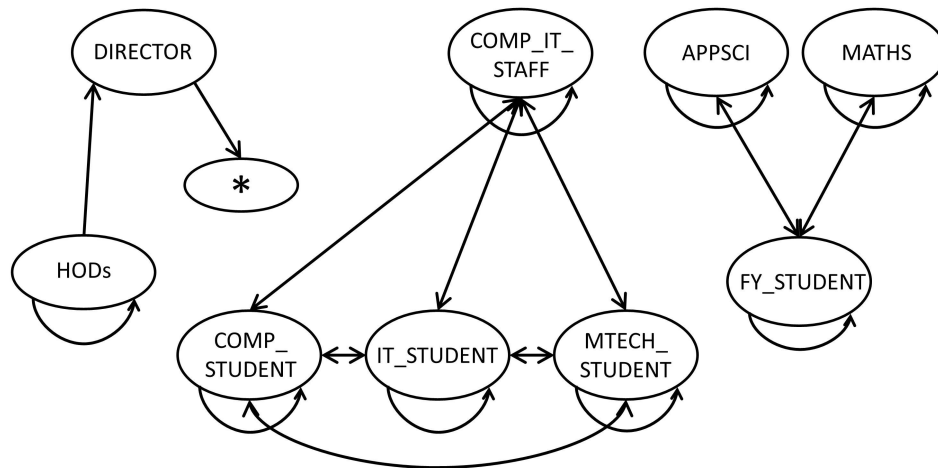


Figure 5.1: Example of a hierarchy related to computer department in an institute. Note: The figure is not representative of any hierarchy and is meant only for illustration purpose to explain the concept of hierarchy. Organizations can decide the format in which they want to represent their hierarchy.

relationship between the sender and recipients. Note the use of the term ‘meaningful’ relationship, and not just relationship. An example of a meaningful relationship is staff of computer department sending emails to students of computer department. Whereas the staff of computer department sending email to metallurgy department student is not meaningful. We discuss more on this after we describe how we establish meaningful relationships.

We introduce the concept of Organization Email Sending Hierarchy. This hierarchy is a graph made of Nodes and Links between them, which shows the meaningful email sending relationships in the organization. The Nodes can be individual users or groups of users, and are connected using Links. Links are directional. A Link from node A to node B means that node A has a meaningful email sending relationship with node B. The same is not applicable vice-versa unless there is a link from Node B to A.

Though the nodes in the hierarchy can be individual users, to reduce the complexity of the hierarchy, the nodes could also be ‘groups’ within the organization (e.g.: Teams, Classes, Departments, etc.). Users who have more responsibility/rights than what the group has (e.g.: Team Leader, HoD, etc.), can be considered as separate nodes alongwith the group node. In such cases, a user would be part of multiple nodes and would have different links with other nodes. Various levels/subsets of groups can be used in the hierarchy. A simple example of hierarchy related to a department in an organization is shown in Figure 5.1.

We call this Email Sending Hierarchy and not Email Sending Patterns, since we do not consider only frequent patterns per se. E.g.: Director of an institute may not be sending emails very frequently to students in the institute and the pattern of sending could be arbitrary, but it is a meaningful relationship. Due to this, we cannot judge anomalous behavior by looking at factors such as how many previous days the sender has sent emails to the recipient. Hence, we do not just use email sending patterns or number of previous days the sender has sent emails to the recipients, but look for meaningful relationships in email sending for construction of hierarchy. Administrators can embed the domain knowledge of their problem space to construct such hierarchy.

Administrators can leverage email history of the organization using learning techniques to find email sending patterns. As mentioned above, we look for ‘meaningful’ patterns and not just frequent patterns. Therefore, administrators can use learning techniques to get a preliminary view of the organization’s patterns, and then remove and add relationships depending on their needs. The hierarchy can be designed easily and saved to database using JavaScript libraries such as GoJS, SigmaJS, VisJS, etc. (or libraries available in any other language of choice). This is a Context-based feature which checks whether it is ‘meaningful’ for a sender to send email to its recipients. There are legitimate exceptions where email sending might not have meaningful relationships, but these are presumably rare.

More features can be added to this category depending on the organization. For example, in our case, we consider an email to be more suspicious if there is no link in the hierarchy for the particular email and the recipients are included in BCC.

5.1.2 IP Feature Category

An account’s compromise can take place from anywhere in the world. For efficiency and ease of use, organizational users are allowed to access their accounts remotely. This also provides attackers with opportunities to infiltrate the network. This feature category checks whether the email is sent from a known geographic location. The check is both network-wide as well as per-user. The first feature in this category is based on sender’s City obtained from Geo-IP mapping. This feature further has two sub-features. These give

- i. the number of times the sender has sent emails from the city, and
- ii. the number of other users that have sent emails from the city.

In both cases, lesser the value, more is the suspiciousness. There might be a case where

the sender has sufficient number of emails from the geolocated city, but there are not enough users who have sent emails from that city. This can happen when a particular user has to visit the city many number of times and requires sending emails from that location. Hence, if the first sub-feature gives a sufficient high value, we disregard the second sub-feature value. Yet, the second sub-feature serves a major purpose, since many students in our organization send emails rarely. So the per-user historical profile is not built sufficiently. In such cases, this network-wide sub-feature plays major role which gives the count of other users who have sent emails from the city. We tested various Geo-IP databases and found that IP2Location gives more accurate city mapping than other services for our case, followed by MaxMind.

IP-City mapping is not always reliable and has inaccuracies in some cases, hence we use a second feature in this category, which is based on the ASN (Autonomous System Number) obtained from the IP-ASN mapping. An ASN can be finer grained than geolocated city, and sometimes vice-versa. But it is a good enough substitute in practice. Similar to the City based sub-features, this feature also has two sub-features. These give

- i. the number of times the sender has sent emails from the ASN, and
- ii. the number of other users that have sent emails from the ASN.

Like the City based sub-features mentioned above, in this case too, if the first feature gives a sufficient high value for the first sub-feature, we disregard the second sub-feature value. For the IP-ASN mapping, we use MaxMind's ASN database.

Both the features in this category are history-based.

A problem arises when a legitimate user sends email during travel. To determine whether the user is travelling, Javed [40] has used ASNs of airports to check whether user is accessing the account from an airport using the public wifi. But this can also be exploited by the attacker if large number of such ASNs corresponding to public places are whitelisted. And in our case there are many possibilities such as public wifi in train, bus, etc. So we decided not to consider the case of travelling of user.

Due to dynamically changing IPs within our organization and also for users accessing accounts remotely, we cannot check whether the user has used the IP address before. A user's machine can be allocated different IP on different occasions within our organization.

5.1.3 Display Name Feature Category

Lateral spear phishing attacks can also be accompanied by display name spoof, as is found in four such attack instances in our dataset. As shown in various previous studies, even technically sound people may fall for display name spoof [4, 6]. For example, a student’s compromised account been given the display name of a lab assistant (Computer lab assistant <studentemailid@coep.ac.in>). The receivers would respond to the email, perceiving that the sender is from the real lab assistant.

Authentication techniques such as SPF/DKIM/DMARC [53] do not work in this case since the email is intra-domain and is being sent from a legitimate account (which is compromised).

The first feature in this category gives the number of times the sender has sent emails with the current display name. This is a history-based feature.

Many students in the institute send emails rarely. Hence, sufficient display name records are not available for such students. The second feature in this category checks whether display name is lexically similar to email address, if there are insufficient number of current display name records for the sender. Many times, organizations set email addresses of employees based on their real name. Hence, lexicographic comparison can work even if employees don’t use exact name (ex. using nicknames). For example, email addresses in our institute have a part of the name (first name and/or surname) of the user. We use this structure to lexicographically compare the email sender’s email address and display name. This is a context-based feature.

We do not consider the trustworthiness/reputation of the spoofed display name, since we have instances where people have fallen for these attacks without the use of authoritative/known display names.

5.1.4 FQDN Feature Category

As seen from Table 4.2, 22 attack instances in our dataset used a phishing URL in the email. URLs can have many variations, such as various parameters in the URL. Instead of using the full URL, we use its Fully Qualified Domain Name (FQDN) to deal with fine granularity and yet not be too coarse. We consider a FQDN which is newer to the organization to be more risky. The first feature checks whether the FQDN is seen for the first time. We notice that most of the FQDNs in our dataset appeared only once in the entire observation period. From Figure 5.2, it can be seen that only 35% of the FQDNs




Figure 5.2: Complementary cumulative distribution function (CCDF) of number of times a FQDN was seen in our dataset.

were seen more than once, or in other words, 65% of the FQDNs were seen only once.

Due to this fact, we cannot use familiarity of FQDNs to our organization alone as a feature. This can also be abused by attackers by sending large number of emails containing a phishing FQDN to increase the familiarity of the FQDN to the organization. We keep track of the *reputation* of the FQDNs *within* the organization. If the FQDN is seen before, then we use our second feature, which checks for the reputation of the FQDN within the organization. We define the reputation of a FQDN as the suspiciousness of the emails it has appeared in.

We do not use domain reputation such as PageRank etc., since adversaries can avoid detection by hosting their pages on reputed domains (E.g.: weebly, wordpress, etc.). This second feature gives the past reputation of the FQDN within the organization. The reputation of a FQDN is dependent on the scores of the emails in which it has been seen. The calculation of the reputation of the FQDN is discussed in the next section when we describe the scoring technique. If a phishing FQDN is seen more number of times, the first feature would fail, but since we maintain the reputation of the FQDN, the second feature would aid us due to the historical low reputation of the FQDN in our system.

Both features in the category are history-based. We do not use rest of the text in the email body. We only use the Fully Qualified Domain Name of any URL that is present in the email body.

We have summarized the features in Table 5.1.

Table 5.1: Summary of features fed into our Scoring system

Feature	Nature
Behavioral Context Feature Category	
Presence of Link between the Nodes in Hierarchy	Context-based
Receivers in BCC	Context-based
IP Feature Category	
City Features	History-based
ASN Features	History-based
Display Name Feature Category	
Display Name History	History-based
Lexicographical Matching	Context-based
FQDN Feature Category	
Familiarity to Organization	History-based
<i>Reputation within</i> Organization	History-based

5.2 LSP Scoring

In the previous section, we described the features we use to characterize individual emails. We now describe how the individual feature categories are aggregated to provide a score for an email and take actions correspondingly. For simplicity, we abbreviate Lateral Spear Phishing as LSP, hence the scoring technique is called as LSP Scoring.

The scoring technique gives a final score in the range of 0.0 to 1.0. Higher score corresponds to email being more suspicious. We consider a score greater than 0.5 to be a lateral spear phishing email. All feature category scores are combined together using simple summation to give the final LSP score for the email. A feature category gets its score from the scores of its individual features. Each feature category has a maximum allowed score such that it does not cross 0.5. This is done to ensure that only one feature category does not get to classify an email as lateral spear phishing. Hence, for an email to be classified as lateral spear phishing, a non-zero score is required from atleast two feature categories. Individual feature scores in the categories are normalized accordingly.

Each history-based feature has its own threshold, decided by the administrators in the organization. If the value given by a history-based feature exceeds its threshold, the feature does not give any score. The score increases linearly with lesser historical values.

Once the LSP score for the email is calculated, the reputation of any FQDNs in the email bodies is re-calculated. This reputation is calculated using the weighted average of the

Algorithm 1: LSP Scoring

```

process_email(email E):
1   $F \leftarrow \text{parse\_fields}(E)$ 
2   $s_1 \leftarrow \text{get\_behav\_score}(F[\text{sender emailid}, \text{receivers}])$ 
3   $s_2 \leftarrow \text{get\_ip\_score}(F[\text{sender ip}])$ 
4   $s_3 \leftarrow \text{get\_dn\_score}(F[\text{sender emailid}, \text{sender displayname}])$ 
5   $s_4 \leftarrow \text{get\_fqdn\_score}(F[\text{fqdns}])$ 
6   $\text{lsp\_score} \leftarrow s_1 + s_2 + s_3 + s_4$ 
   if  $\text{lsp\_score} > 0.5$  then
7     flag email
8     update tables except history-based tables
   else
9     update all tables
10 return  $\text{lsp\_score}$ 

```

LSP score of the current email and past reputation of the FQDN, with the number of previous days the FQDN has been received. This new reputation is used the next time this FQDN is seen.

Features such as IP, display name require prior history to build historical profiles and give accurate results. A bootstrapping period is required for these history-based features. Although more history would increase the chances of the system to give more accurate results, it also means that the system would require large amount of history to give good results. A tradeoff needs to be obtained in this case such that sufficient amount of history is used, and yet it is not too much. Considering this, our training dataset contains initial 4 months of emails.

In this bootstrap period, though IP and display name history will be built, the FQDN reputations of URLs in the emails also need to be determined. FQDN reputations are derived from LSP scores of the emails. During the bootstrap period, if we simply determine LSP scores of emails directly using all the features without the pre-existence of historical profiles, we get large number of false positives. This is mainly due to the lack of history for IP and display name features. Also, benign FQDNs get bad reputation if they are present in these earlier emails and further contribute to the false positives. To overcome this issue of building history (for IP, display name features) as well as finding LSP scores of the same emails (in order to calculate FQDN reputations), we use two steps in the training phase. In the first step, we use the initial four months of emails to build IP and display name history. Then in the second step, we use the same four months of emails

to find LSP scores of the emails using all features (IP and display names tables are only referred to and not updated in this step, since they are filled with data of same emails in first step). As the LSP scores of emails are determined, FQDN reputations of any URLs in the email bodies are updated in the database. Once the database is filled with IP, display name and FQDN reputation data from training emails, we run the algorithm on the emails after the bootstrap period.

There were inaccuracies even in the best performing Geo-IP database in our case. Due to this, City feature in the IP Feature Category gives a non-zero score owing to the absence of sufficient history of the incorrect City. There were considerably lesser number of inaccuracies in the IP-ASN database. Hence, if City feature gives a non-zero score due to insufficient history, but ASN feature gives a zero score due to sufficient history, we consider the IP Feature Category as a whole to have zero score. We assume that if the system contains sufficient historical profile for an ASN, the system should also contain sufficient historical profile for the corresponding City.

An attacker may try to contaminate our historically built profiles by sending large number of emails from an anomalous IP and/or using spoofed display name. Therefore, if an email gets a score greater than 0.5, we do not add its history based values to our database. Algorithm 1 shows the procedure of LSP Scoring in short.

Chapter 6

Evaluation

We evaluate our scoring technique on the dataset which we collected from 40 volunteers in COEP, amounting to 3.5 years' worth of emails. The dataset contains 27 attack instances which correspond to 323 lateral spear phishing emails. Using our scoring system, we were able to detect 22 attack instances and 1 attack instance partially (some emails in the attack instance were detected), which correspond to 280 emails. Additionally, out of the 27 attack instances, 25 instances were known previously. We detected 2 attack instances in the dataset which were not known to us earlier. Out of these two, one attack instance was highly targeted and corresponds to only one lateral spear phishing email. We obtained false positive rate of 0.88% with true positive rate of 86.69%. The overall accuracy is 98.79%. Results of various performance metrics are shown in Table 6.1.

The most common reason of false positives is incorrect City result given by the Geo-IP database. On manual check, we found that the senders were actually sending from the same city as the organization, but the Geo-IP database had incorrect IP-City mapping. This error alongwith smaller values given by other feature categories led to the LSP score cross 0.5. Although none of the benign emails got a score of above 0.65.

Figure 6.1 shows how the LSP Scoring system scores benign and lateral spear phishing emails. 99.1% of benign emails have score below 0.5. And 86.6% of lateral spear phishing emails score over 0.54. This indicates that the compact feature set provides enough separation for correctly identifying benign and lateral spear phishing emails.

Table 6.1: Comparison of performance of our Scoring Technique with that of Anomaly Detection Techniques

Algorithm	True Positive	False Negative	True Negative	False Positive	Accuracy	F1 Score
LSP Scoring	86.69%	13.31%	99.12%	0.88%	98.79%	79.32%
Local Outlier Factor	96.28%	3.72%	2.22%	97.78%	4.74%	5.14%
One-Class SVM	99.07%	0.93%	45.32%	54.68%	46.76%	9.07%
Robust Covariance	91.95%	8.05%	96.3%	3.7%	96.18%	56.3%

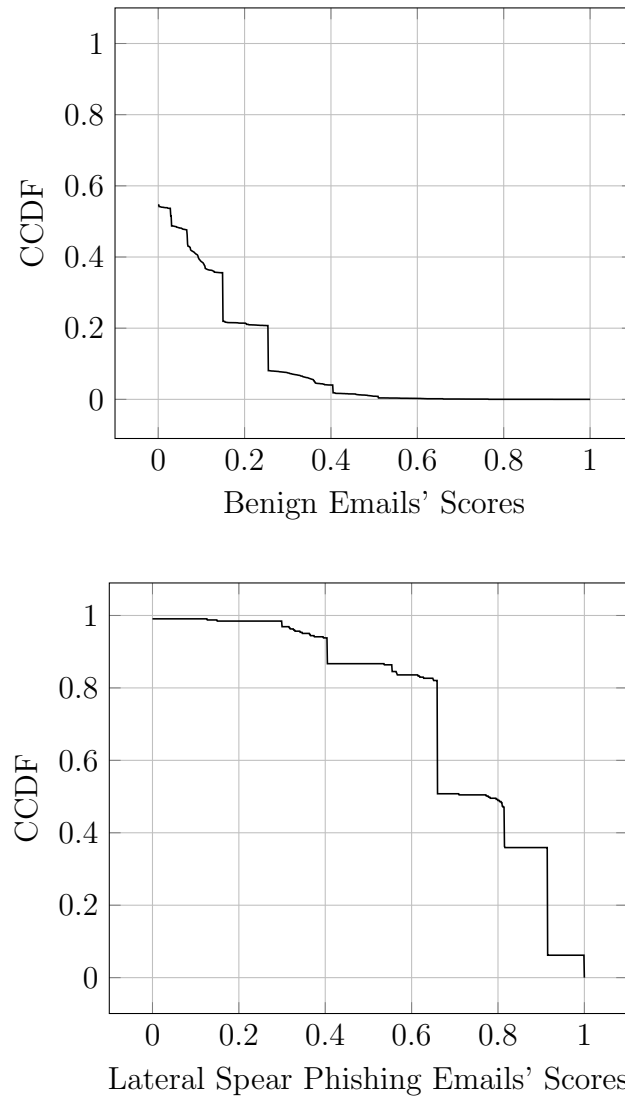


Figure 6.1: CCDF of LSP scores for benign and lateral spear phishing emails given by the LSP Scoring system.

6.1 Comparison with Anomaly Detection Techniques

We compare our system's performance with that of the anomaly detection techniques based on machine learning from literature [12, 31], viz. Local Outlier Factor, One-Class SVM, and Robust Covariance. For each email, we constructed a feature vector from the scores given by each of the feature categories and normalizing them. We have shown the comparison of the results in Table 6.1. Although One-Class SVM gives the highest true positive rate, it also gives a very unacceptable false positive rate. Similar results can be observed in other techniques. All the compared anomaly detection techniques have false positive rates impractical for use in detection of such attacks. Our False Positive Rate,

True Negative Rate, Accuracy and F1 Score are much better than other algorithms. Even though the True Positive and False Negative Rates of other algorithms are better than ours, it is mainly due to our effort to keep False Positive Rate to the minimum which is a very important factor for the use of a system in practice.

Standard anomaly detection algorithms suffer from the limitation of determining events as anomalous even if one of their feature values are anomalous. Due to the diversity of benign behavior found in organizations, various benign instances can have anomalous values for different features. Since our algorithm incorporates domain knowledge and does not determine an event as anomalous merely due to high value given by one feature, it fares better than standard algorithms. These algorithms also suffer from the limitation of either assuming data to be of a particular distribution, or need parameters to be set by the administrator to obtain acceptable results. The distribution of specialized attacks such as these are unknown and can vary from organization to organization. Setting the parameters requires fine tuning by the administrator and there is no established way of selecting the optimal parameter values.

Chapter 7

Discussion

Cyber space is in a continuous state of tug of war between attackers and defenders. Our motive behind this research is to make it difficult for attackers to conduct successful lateral spear phishing attacks. Our scoring system uses context-based and history-based features to detect lateral spear phishing emails. The system matures and continuously refines itself as it processes more number of emails. Administrators can take action in real-time and stop attack from causing further damage by revoking access of compromised account. As shown in the previous chapter, we are able to obtain high accuracy with the compact feature set. Additionally, due to the compact feature set, the decision boundary is less complex. Depending on the problem space, administrators can add required features to the system. Since these attacks are highly specialized, domain knowledge can help to a large extent in detecting them. As we do not update the history based values if an email has been marked as lateral spear phishing, this would not allow the attacker to easily contaminate the historically built profiles.

7.1 Limitations

An obvious limitation of the system is that it does not detect spear phishing attacks from email addresses outside the organization. Since our system does not consider spear phishing emails coming from email addresses outside COEP, such attacks would be ignored by our system.

To build historical profiles, the system needs prior training. We have given the detailed results in previous chapter with the training period of 4 months. We also tested the system with lesser training periods. These tests gave higher false positive rates than the test with 4 month training period. This shows that a minimum required amount of

training is needed to give better results. This period depends on the density of the users' activity.

From the email bodies, we consider only FQDN of any present URLs, and not any textual content. This decision was mainly taken due to privacy and confidentiality reasons. A detector may use the textual content to build stylometric models and use NLP techniques.

As explained in 5.1.2, we do not consider the case of travelling user. This would result in the system giving inaccurate score in the IP Feature Category which would change the total LSP score and hence cause misclassification in such cases.

7.2 Chrome Extension - COEP Kumpan

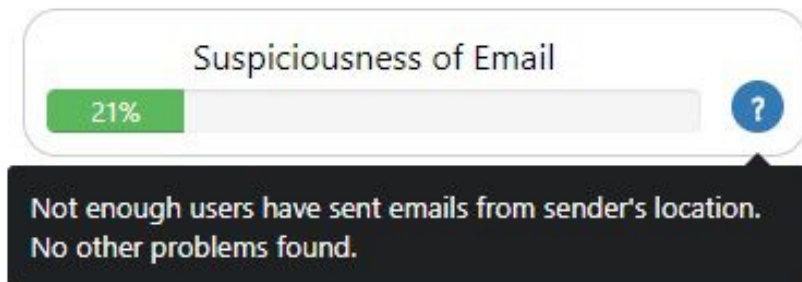
The proposed technique is primarily aimed to be implemented at server side by email providers who provide email services to organizations, or by organizations that manage their own email server. However, this is not always possible. Depending on the organization's requirements, the organization may decide to allow all intra-domain emails to reach the recipients and let the end user decide what action to take on the emails. Or the organization might be using an email service which does not support such detection mechanism. In such cases, to aid users in their decision while dealing with an email, organizations can implement a technique such as ours, within their organization with the help of plugins.

We developed a Chrome Extension, called COEP Kumpan, which displays a UI element on Outlook Web (used by majority of COEP users) to guide users in their decision making. The extension uses an API we provided for our scoring system. Kumpan, in Marathi language, means fence - in our case, to safeguard the organization.


The extension is developed from the point of view of preventative defense to provide users with warning messages. The design of the UI element displayed by the extension is aimed to nudge with relevant information of the email [54], in order to capture users' attention and guide them to take informed decision. Figure 7.1 shows screenshots of the UI element shown to user. In addition to the chrome extension, we developed the UIs necessary for administrators in COEP to use the system, such as constructing the hierarchy and viewing details of detected attacks. Since we do not work on detecting spear phishing emails coming from email addresses not belonging to COEP domain, we show a generic warning message for all emails not from COEP domain, and give instructions to user to proceed carefully for such emails.



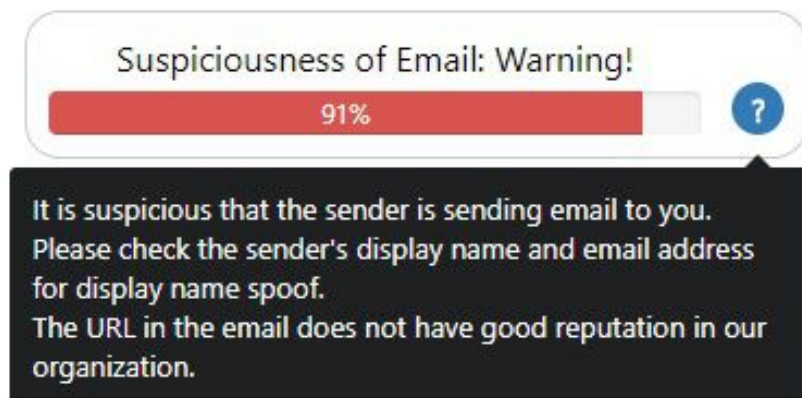
(a)



(b)



(c)



(d)

Figure 7.1: (a) shows an example of the UI element in an email in Outlook Web. (b)-(d) show examples of UI element with messages, either nudged or on mouse hover based on user preferences.

Chapter 8

Conclusion

In this dissertation, we studied a specialized class of social engineering attack - Lateral Spear Phishing attack. Spear phishing emails in such attacks are particularly hard to detect since they are sent from a compromised account within the organization. Our aim was to design a practical, deployable and real-time system to detect such attacks. We collected 3.5 years' worth of emails from 40 volunteers in our organization and provided the characterization of this non-labelled dataset. We derived various feature categories based on our observation of the attacks on the organization. These features were fed into our proposed Scoring technique to detect Lateral Spear Phishing emails. We evaluated our scoring technique on our real-world email dataset. The results show that the scoring technique achieved an accuracy of 98.79% and false positive rate of 0.88%. Also, we were able to detect 2 attack instances which were not known to us earlier. We compared the performance of our scoring technique with that of machine learning based anomaly detection techniques from literature, and show that our scoring technique gives better overall results. We also developed a Chrome extension - COEP Kumpan - for use by COEP to help its users in taking better informed decisions. Domain knowledge can play a crucial role in detection of such specialized attacks. Since organizations currently depend mostly on users for reporting spear phishing attacks, our research can serve as an improvement in detecting fair number of true positives, keeping low false positive rate, and also identifying previously unknown attacks.

Future Work. As part of future work, we aim to evaluate the chrome extension through users' survey, and gauge its effect on decision making of various users in COEP. We would also look into further improving the results and engineering of more features to be incorporated into the system. Also, given the availability of data, we could explore and combine identifying attackers during login stage with the post-login email activity we explored in this work.

Bibliography

- [1] Agari. Agari Global DMARC Adoption Report, 2017.
- [2] Syed Mubashir Ali. Integration of information security essential controls into information technology infrastructure library-A proposed framework. *International Journal of Applied Science and Technology*, 4(1), 2014.
- [3] A. Almomani, B. B. Gupta, S. Atawneh, A. Meulenberg, and E. Almomani. A Survey of Phishing Email Filtering Techniques. *IEEE Communications Surveys Tutorials*, 15(4):2070–2090, 2013.
- [4] Mohamed Alsharnouby, Furkan Alaca, and Sonia Chiasson. Why phishing still works: User strategies for combating phishing attacks. *International Journal of Human-Computer Studies*, 82:69–82, October 2015.
- [5] Calvin Ardi and John Heidemann. AuntieTuna: Personalized Content-Based Phishing Detection. In *Proceedings of the NDSS Workshop on Usable Security*, San Diego, California, USA, February 2016. The Internet Society.
- [6] Zinaida Benenson, Freya Gassmann, and Robert Landwirth. Unpacking Spear Phishing Susceptibility. In *Financial Cryptography and Data Security*, Lecture Notes in Computer Science, pages 610–627. Springer, Cham, April 2017.
- [7] Russell Bandom. Two-factor authentication is a mess, July 2017.
- [8] Cristian Bravo-Lillo, Saranga Komanduri, Lorrie Faith Cranor, Robert W. Reeder, Manya Sleeper, Julie Downs, and Stuart Schechter. Your Attention Please: Designing Security-decision UIs to Make Genuine Risks Harder to Ignore. In *Proceedings of the Ninth Symposium on Usable Privacy and Security*, SOUPS '13, pages 6:1–6:12, New York, NY, USA, 2013. ACM.
- [9] Casey Inez Canfield, Baruch Fischhoff, and Alex Davis. Quantifying Phishing Susceptibility for Detection and Behavior Decisions. *Human Factors*, 58(8):1158–1172, December 2016.

-
- [10] Ye Cao, Weili Han, and Yueran Le. Anti-phishing Based on Automated Individual White-list. In *Proceedings of the 4th ACM Workshop on Digital Identity Management*, DIM '08, pages 51–60, New York, NY, USA, 2008. ACM.
 - [11] D. D. Caputo, S. L. Pfleeger, J. D. Freeman, and M. E. Johnson. Going Spear Phishing: Exploring Embedded Training and Awareness. *IEEE Security Privacy*, 12(1):28–38, January 2014.
 - [12] Varun Chandola, Arindam Banerjee, and Vipin Kumar. Anomaly detection: A survey. *ACM computing surveys (CSUR)*, 41(3):15, 2009.
 - [13] Neil Chou, Robert Ledesma, Yuka Teraguchi, and John C. Mitchell. Client-Side Defense Against Web-Based Identity Theft. In *NDSS*, 2004.
 - [14] Dan Conway, Ronnie Taib, Mitch Harris, Kun Yu, Shlomo Berkovsky, and Fang Chen. A Qualitative Investigation of Bank Employee Experiences of Information Security and Phishing. In *Thirteenth Symposium on Usable Privacy and Security (SOUPS 2017)*, pages 115–129. USENIX Association, 2017.
 - [15] Qian Cui, Guy-Vincent Jourdan, Gregor V. Bochmann, Russell Couturier, and Iosif-Viorel Onut. Tracking Phishing Attacks Over Time. In *Proceedings of the 26th International Conference on World Wide Web*, WWW '17, pages 667–676, Republic and Canton of Geneva, Switzerland, 2017. International World Wide Web Conferences Steering Committee.
 - [16] Tamara Denning, Adam Lerner, Adam Shostack, and Tadayoshi Kohno. Control-Alt-Hack: The Design and Evaluation of a Card Game for Computer Security Awareness and Education. In *Proceedings of the 2013 ACM SIGSAC Conference on Computer & Communications Security*, CCS '13, pages 915–928, New York, NY, USA, 2013. ACM.
 - [17] P. Dewan, A. Kashyap, and P. Kumaraguru. Analyzing social and stylometric features to identify spear phishing emails. In *2014 APWG Symposium on Electronic Crime Research (eCrime)*, pages 1–13, September 2014.
 - [18] Rachna Dhamija, J. D. Tygar, and Marti Hearst. Why Phishing Works. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, CHI '06, pages 581–590, New York, NY, USA, 2006. ACM.
 - [19] DNSBL. DNSBL, September 2017. Page Version ID: 800548089.
 - [20] Julie S. Downs, Mandy Holbrook, and Lorrie Faith Cranor. Behavioral Response to Phishing Risk. In *Proceedings of the Anti-phishing Working Groups 2Nd Annual*

- eCrime Researchers Summit*, eCrime '07, pages 37–44, New York, NY, USA, 2007. ACM.
- [21] Julie S. Downs, Mandy B. Holbrook, and Lorrie Faith Cranor. Decision Strategies and Susceptibility to Phishing. In *Proceedings of the Second Symposium on Usable Privacy and Security*, SOUPS '06, pages 79–90, New York, NY, USA, 2006. ACM.
- [22] S. Duman, K. Kalkan-Cakmakci, M. Egele, W. Robertson, and E. Kirda. EmailProfiler: Spearphishing Filtering with Header and Stylometric Features of Emails. In *2016 IEEE 40th Annual Computer Software and Applications Conference (COMPSAC)*, volume 1, pages 408–416, June 2016.
- [23] Manuel Egele, Gianluca Stringhini, Christopher Krügel, and Giovanni Vigna. COMPA: Detecting Compromised Accounts on Social Networks. In *NDSS*, 2013.
- [24] Serge Egelman, Lorrie Faith Cranor, and Jason Hong. You’ve Been Warned: An Empirical Study of the Effectiveness of Web Browser Phishing Warnings. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, CHI '08, pages 1065–1074, New York, NY, USA, 2008. ACM.
- [25] Adrienne Porter Felt, Robert W. Reeder, Alex Ainslie, Helen Harris, Max Walker, Christopher Thompson, Mustafa Emre Acer, Elisabeth Morant, and Sunny Consolvo. Rethinking Connection Security Indicators. In *Twelfth Symposium on Usable Privacy and Security (SOUPS 2016)*, pages 1–14, Denver, CO, 2016. USENIX Association.
- [26] Jon Fingas. Florida phishing attack exposes data for 30,000 Medicaid recipients, January 2018.
- [27] Fireeye. Best Defense Against Spear-Phishing Attacks, January 2018.
- [28] Edwin Donald Frauenstein and Rossouw von Solms. An Enterprise Anti-phishing Framework. In *Information Assurance and Security Education and Training*, IFIP Advances in Information and Communication Technology, pages 196–203. Springer, Berlin, Heidelberg, July 2009.
- [29] David Mandell Freeman, Sakshi Jain, Markus Dürmuth, Battista Biggio, and Giorgio Giacinto. Who Are You? A Statistical Approach to Measuring User Authenticity. In *NDSS*, 2016.
- [30] Gerry Gaffney. The myth of the stupid user | Information & Design, 2011.

-
- [31] Markus Goldstein and Seiichi Uchida. A Comparative Evaluation of Unsupervised Anomaly Detection Algorithms for Multivariate Data. *PLOS ONE*, 11(4):e0152173, April 2016.
 - [32] B. B. Gupta, Aakanksha Tewari, Ankit Kumar Jain, and Dharma P. Agrawal. Fighting against phishing attacks: state of the art and future challenges. *Neural Computing and Applications*, 28(12):3629–3654, December 2017.
 - [33] Mingxing He, Shi-Jinn Horng, Pingzhi Fan, Muhammad Khurram Khan, Ray-Shine Run, Jui-Lin Lai, Rong-Jian Chen, and Adi Sutanto. An Efficient Phishing Webpage Detector. *Expert Syst. Appl.*, 38(10):12018–12027, September 2011.
 - [34] Ryan Heartfield and George Loukas. A Taxonomy of Attacks and a Survey of Defence Mechanisms for Semantic Social Engineering Attacks. *ACM Comput. Surv.*, 48(3):37:1–37:39, December 2015.
 - [35] Grant Ho, Aashish Sharma, Mobin Javed, Vern Paxson, and David Wagner. Detecting Credential Spearphishing in Enterprise Settings. In *26th USENIX Security Symposium (USENIX Security 17)*, pages 469–485, Vancouver, BC, 2017. USENIX Association.
 - [36] Jason Hong. The State of Phishing Attacks. *Commun. ACM*, 55(1):74–81, January 2012.
 - [37] Imperva. Phishing made easy: Time to rethink your prevention strategy?, December 2016.
 - [38] Tom N. Jagatic, Nathaniel A. Johnson, Markus Jakobsson, and Filippo Menczer. Social Phishing. *Commun. ACM*, 50(10):94–100, October 2007.
 - [39] Ankit Kumar Jain and B. B. Gupta. Phishing Detection: Analysis of Visual Similarity Based Approaches. *Security and Communication Networks*, 2017, 2017.
 - [40] Mobin Javed. *Detecting Credential Compromise in Enterprise Networks*. PhD thesis, Electrical Engineering and Computer Sciences University of California at Berkeley, December 2016.
 - [41] M. Khonji, Y. Iraqi, and A. Jones. Mitigation of spear phishing attacks: A Content-based Authorship Identification framework. In *2011 International Conference for Internet Technology and Secured Transactions*, pages 416–421, December 2011.
 - [42] Brian Krebs. Equifax or Equiphish? — Krebs on Security, September 2017.

-
- [43] Brian Krebs. The Market for Stolen Account Credentials — Krebs on Security, December 2017.
- [44] Ponnurangam Kumaraguru. *Phishguru: A System for Educating Users About Semantic Attacks*. PhD thesis, Carnegie Mellon University, Pittsburgh, PA, USA, 2009.
- [45] Ponnurangam Kumaraguru, Justin Cranshaw, Alessandro Acquisti, Lorrie Cranor, Jason Hong, Mary Ann Blair, and Theodore Pham. School of Phish: A Real-world Evaluation of Anti-phishing Training. In *Proceedings of the 5th Symposium on Usable Privacy and Security*, SOUPS '09, pages 3:1–3:12, New York, NY, USA, 2009. ACM.
- [46] Ponnurangam Kumaraguru, Steve Sheng, Alessandro Acquisti, Lorrie Faith Cranor, and Jason Hong. Teaching Johnny not to fall for phish. *ACM Transactions on Internet Technology (TOIT)*, 10(2):7, 2010.
- [47] Kaspersky Lab. Phishing for cryptocurrencies: How bitcoins are stolen, January 2018.
- [48] Comodo Threat Research Labs. Phishing got Darker and Smarter, January 2018.
- [49] Aron Laszka, Jian Lou, and Yevgeniy Vorobeychik. Multi-defender Strategic Filtering Against Spear-phishing Attacks. In *Proceedings of the Thirtieth AAAI Conference on Artificial Intelligence*, AAAI'16, pages 537–543, Phoenix, Arizona, 2016. AAAI Press.
- [50] Justin Ma, Lawrence K. Saul, Stefan Savage, and Geoffrey M. Voelker. Beyond Blacklists: Learning to Detect Malicious Web Sites from Suspicious URLs. In *Proceedings of the 15th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, KDD '09, pages 1245–1254, New York, NY, USA, 2009. ACM.
- [51] Aaron Mak. Oh Great, a Hacking Group Linked to North Korea Is Getting Very Good at Targeting Bitcoin Owners, February 2018.
- [52] Sarah Meyer. Phishing Attacks: Insights from More Than 1,000 Free Phishing Kits - Page 2 of 2, January 2018.
- [53] Rob Mueller. SPF, DKIM & DMARC: email anti-spoofing technology history and future, December 2016.
- [54] James Nicholson, Lynne Coventry, and Pam Briggs. Can we fight social engineering attacks by social means? Assessing social salience as a means to improve phish detection. In *Thirteenth Symposium on Usable Privacy and Security (SOUPS 2017)*, pages 285–298. USENIX Association, 2017.

-
- [55] National Institute of Standards and Technology. Framework for Improving Critical Infrastructure Cybersecurity.
 - [56] Daniela Oliveira, Harold Rocha, Huizi Yang, Donovan Ellis, Sandeep Dommaraju, Melis Muradoglu, Devon Weir, Adam Soliman, Tian Lin, and Natalie Ebner. Dissecting Spear Phishing Emails for Older vs Young Adults: On the Interplay of Weapons of Influence and Life Domains in Predicting Susceptibility to Phishing. In *Proceedings of the 2017 CHI Conference on Human Factors in Computing Systems*, CHI '17, pages 6412–6424, New York, NY, USA, 2017. ACM.
 - [57] Daniela Oliveira, Harold Rocha, Huizi Yang, Donovan Ellis, Sandeep Dommaraju, Melis Muradoglu, Devon Weir, Adam Soliman, Tian Lin, and Natalie Ebner. Dissecting Spear Phishing Emails for Older vs Young Adults: On the Interplay of Weapons of Influence and Life Domains in Predicting Susceptibility to Phishing. In *Proceedings of the 2017 CHI Conference on Human Factors in Computing Systems*, CHI '17, pages 6412–6424, New York, NY, USA, 2017. ACM.
 - [58] Exchange Online. Configure your spam filter policies: Exchange Online Protection Help, December 2017.
 - [59] A. Pecchia, A. Sharma, Z. Kalbarczyk, D. Cotroneo, and R. K. Iyer. Identifying Compromised Users in Shared Computing Infrastructures: A Data-Driven Bayesian Network Approach. In *2011 IEEE 30th International Symposium on Reliable Distributed Systems*, pages 127–136, October 2011.
 - [60] PhishTank. PhishTank > Frequently Asked Questions (FAQ), 2017.
 - [61] PhishTank. PhishTank > Friends of PhishTank, 2017.
 - [62] P. Prakash, M. Kumar, R.R. Kompella, and M. Gupta. PhishNet: Predictive Blacklisting to Detect Phishing Attacks. In *2010 Proceedings IEEE INFOCOM*, pages 1–5, March 2010.
 - [63] Gowtham Ramesh, Ilango Krishnamurthi, and K. Sampath Sree Kumar. An efficacious method for detecting phishing webpages through target domain identification. *Decision Support Systems*, 61(Supplement C):12–22, May 2014.
 - [64] Dylan Sachs. How to Take Down a Phishing Site: 5 Crucial Steps, 2013.
 - [65] Stuart E. Schechter, Rachna Dhamija, Andy Ozment, and Ian Fischer. The Emperor’s New Security Indicators. In *Proceedings of the 2007 IEEE Symposium on Security and Privacy*, SP '07, pages 51–65, Washington, DC, USA, 2007. IEEE Computer Society.

-
- [66] Steve Sheng, Mandy Holbrook, Ponnurangam Kumaraguru, Lorrie Faith Cranor, and Julie Downs. Who Falls for Phish?: A Demographic Analysis of Phishing Susceptibility and Effectiveness of Interventions. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, CHI '10, pages 373–382, New York, NY, USA, 2010. ACM.
 - [67] Steve Sheng, Brad Wardman, Gary Warner, Lorrie Faith Cranor, Jason Hong, and Chengshan Zhang. An empirical analysis of phishing blacklists. In *Proceedings of Sixth Conference on Email and Anti-Spam (CEAS)*, 2009.
 - [68] InformationSecurity StackExchange. email - Why does Outlook not block spam sent by employees? - Information Security Stack Exchange, December 2017.
 - [69] Steve Sheng, Bryant Magnien, Ponnurangam Kumaraguru, Alessandro Acquisti, Lorrie Faith Cranor, Jason Hong, and Elizabeth Nunge. Anti-Phishing Phil: The Design and Evaluation of a Game That Teaches People Not to Fall for Phish. In *Proceedings of the 3rd Symposium on Usable Privacy and Security*, SOUPS '07, pages 88–99, New York, NY, USA, 2007. ACM.
 - [70] Gianluca Stringhini and Olivier Thonnard. That Ain't You: Blocking Spearphishing Through Behavioral Modelling. In *Detection of Intrusions and Malware, and Vulnerability Assessment*, Lecture Notes in Computer Science, pages 78–97. Springer, Cham, July 2015.
 - [71] Kurt Thomas, Frank Li, Chris Grier, and Vern Paxson. Consequences of Connectivity: Characterizing Account Hijacking on Twitter. In *ACM Conference on Computer and Communications Security*, 2014.
 - [72] Arun Vishwanath. Habitual Facebook use and its impact on getting deceived on social media. *Journal of Computer-Mediated Communication*, 20(1):83–98, 2015.
 - [73] Jingguo Wang, Yuan Li, and H. Raghav Rao. Overconfidence in Phishing Email Detection. *Journal of the Association for Information Systems*, 17(11), November 2016.
 - [74] Zikai Alex Wen, Yiming Li, Reid Wade, Jeffrey Huang, and Amy Wang. What.Hack: Learn Phishing Email Defence the Fun Way. In *Proceedings of the 2017 CHI Conference Extended Abstracts on Human Factors in Computing Systems*, CHI EA '17, pages 234–237, New York, NY, USA, 2017. ACM.
 - [75] L. Wenying, G. Liu, B. Qiu, and X. Quan. Antiphishing through Phishing Target Discovery. *IEEE Internet Computing*, 16(2):52–61, March 2012.

-
- [76] Colin Whittaker, Brian Ryner, and Marria Nazif. Large-Scale Automatic Classification of Phishing Pages. In *NDSS*, volume 10, page 2010, 2010.
 - [77] Wikipedia. Google Safe Browsing - Wikipedia, 2017.
 - [78] Wombat. Email Security or Anti-Phishing PhyllisTM | Wombat Security, 2017.
 - [79] Min Wu, Robert C. Miller, and Simson L. Garfinkel. Do Security Toolbars Actually Prevent Phishing Attacks? In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, CHI '06, pages 601–610, New York, NY, USA, 2006. ACM.
 - [80] Jing Zhang, R. Berthier, W. Rhee, M. Bailey, P. Pal, F. Jahanian, and W. H. Sanders. Safeguarding academic accounts and resources with the University Credential Abuse Auditing System. In *IEEE/IFIP International Conference on Dependable Systems and Networks (DSN 2012)*, pages 1–8, June 2012.
 - [81] Yue Zhang, Jason I. Hong, and Lorrie F. Cranor. Cantina: A Content-based Approach to Detecting Phishing Web Sites. In *Proceedings of the 16th International Conference on World Wide Web*, WWW '07, pages 639–648, New York, NY, USA, 2007. ACM.

List of Publications

1. A. Bhadane, S. Mane, “Detecting Lateral Spear Phishing Attacks in Organizations”, in: *IET Information Security*. Volume No. 13, Issue No. 2, March 2019, p. 133 140. ISSN Print: 1751-8709, ISSN Online: 1751-8717, DOI: 10.1049/jet-ifs.2018.5090
2. A. Bhadane, S. Mane, “State of Research on Phishing and Recent Trends of Attacks”, in: *i-manager’s Journal on Computer Science*. Volume No. 5, Issue No. 4, December-February 2018, pp. 14-35, ISSN Print: 2347-2227, ISSN Online: 2347-6141, DOI: 10.26634/jcom.5.4.14608
3. A. Bhadane, S. Mane, “State of Research on User Psychology involved in Phishing Attacks”, in: *Computer Society of India (CSI) Communications*, Volume No. 42, Issue No. 2, May 2018, pp. 34-36,39. ISSN: 0970-647X
4. A. Bhadane, S. Mane, “State of Research on User Training against Phishing with Recent Trends of Attacks”, in: *Centre for Advanced Strategic Studies (CASS) Journal*, April-June 2018, pp.38-53, ISSN 2347-9191